

CSCI 2021: x86-64 Control Flow

Chris Kauffman

*Last Updated:
Fri 18 Oct 2019 11:40:36 AM CDT*

Logistics

Reading Bryant/O'Hallaron

- ▶ Ch 3.6: Control Flow
- ▶ Ch 3.7: Procedure calls

Goals

- ▶ Finish Assembly Basics
- ▶ Jumps and Control flow
x86-64
- ▶ Procedure calls

Assignment 3

- ▶ Problem 1: Battery Meter
Assembly Functions (50%)
- ▶ Problem 2: Binary Bomb via
debugger (50%)

Date	Event
Mon 10/14	Asm Control Flow A3 Released
Wed 10/16	Discuss A3
Fri 10/18	Asm Control Wrap
Mon 10/21	Asm Wrap-up
Wed 10/23	Review
Fri 10/25	Exam 2
Mon 10/28	CPU Architecture
Tue 10/29	A3 Due

Control Flow in Assembly and the Instruction Pointer

- ▶ No high-level conditional or looping constructs in assembly
- ▶ Only `%rip`: Instruction Pointer or “Program Counter”: memory address of the next instruction to execute
- ▶ Don’t mess with `%rip` by hand: automatically increases as instructions execute so the next valid instruction is referenced
- ▶ Jump instructions modify `%rip` to go elsewhere
- ▶ Typically label assembly code with positions of instructions that will be the target of jumps
- ▶ **Unconditional Jump** Instructions always jump to a new location.
- ▶ **Comparison / Test** Instruction, sets EFLAGS bits indicating relation between registers/values
- ▶ **Conditional Jump** Instruction, jumps to a new location if certain bits of EFLAGS are set, ignored if bits not set

Examine: Loop Sum with Instruction Pointer (rip)

- ▶ Can see direct effects on rip in disassembled code
- ▶ rip increases corresponding to instruction length
- ▶ Jumps include address for next rip

```
// C Code equivalent
int sum=0, i=1, lim=100;
while(i<=lim){
    sum += i;
    i++;
}
return sum;
```

000000000000005fa <main>:

ADDR	HEX-OPCODES	ASSEMBLY	EFFECT ON RIP
5fa:	48 c7 c0 00 00 00 00	mov \$0x0,%rax	# rip = 5fa -> 601
601:	48 c7 c1 01 00 00 00	mov \$0x1,%rcx	# rip = 601 -> 608
608:	48 c7 c2 64 00 00 00	mov \$0x64,%rdx	# rip = 608 -> 60f
0000000000000060f <LOOP>:			
60f:	48 39 d1	cmp %rdx,%rcx	# rip = 60f -> 612
612:	7f 08	jg 61c <END>	# rip = 612 -> 614 OR 61c
614:	48 01 c8	add %rcx,%rax	# rip = 614 -> 617
617:	48 ff c1	inc %rcx	# rip = 617 -> 61a
61a:	eb f3	jmp 60f <LOOP>	# rip = 61a -> 60f
0000000000000061c <END>:			
61c:	c3	retq	# rip 61c -> return address

Disassembling Binaries

- ▶ Binaries hard to read on their own
- ▶ Many tools exist to work with them, notably `objdump` on Unix
- ▶ Can **disassemble** binary: show “readable” version of contents

```
> gcc -Og loop.s                                # COMPILE AND ASSEMBLE

> file a.out
a.out: ELF 64-bit LSB pie executable, x86-64, version 1 (SYSV),

> objdump -d a.out                             # DISASSEMBLE BINARY
a.out:      file format elf64-x86-64
...
Disassembly of section .text:
...
0000000000000119 <main>:
    1119:      48 c7 c0 00 00 00 00    mov     $0x0,%rax
    1120:      48 c7 c1 01 00 00 00    mov     $0x1,%rcx
    1127:      48 c7 c2 64 00 00 00    mov     $0x64,%rdx
000000000000012e <LOOP>:
    112e:      48 39 d1                  cmp     %rdx,%rcx
    1131:      7f 08                    jg      113b <END>
    1133:      48 01 c8                  add     %rcx,%rax
    1136:      48 ff c1                  inc     %rcx
    1139:      eb f3                    jmp     112e <LOOP>
000000000000013b <END>:
    113b:      c3                      retq
```

FLAGS: Condition Codes Register

- ▶ Most CPUs have a special register with “flags” for various conditions
- ▶ In x86-64 this register goes by the following names

Name	Width	Notes
FLAGS	16-bit	Most important bits in first 16
EFLAGS	32-bit	Name shown in gdb
RFLAGS	64-bit	Not used normally

- ▶ Bits in FLAGS register are **automatically** set based on results of other operations
- ▶ Pertinent examples with conditional execution

Bit	Abbrev	Name	Description
0	CF	Carry flag	Set if last op caused unsigned overflow
6	ZF	Zero flag	Set if last op yielded a 0 result
7	SF	Sign flag	Set if last op yielded a negative
8	TF	Trap flag	Used by gdb to stop after one ASM instruction
9	IF	Interrupt flag	1: handle hardware interrupts, 0: ignore them
11	OF	Overflow flag	Set if last op caused signed overflow/underflow

Comparisons and Tests

Set the EFLAGS register by using comparison instructions

Name	Instruction	Examples	Notes
Compare	cmpX B, A	cmpl \$1,%eax	Like if(eax > 1){...}
	Like: A - B	cmpq %rsi,%rdi	Like if(rdi > rsi){...}
Test	testX B, A	testq %rcx,%rdx	Like if(rdx & rcx){...}
	Like: A & B	testl %rax,%rax	Like if(rax){...}

- ▶ immediates like \$2 must be the first argument B
- ▶ B,A are NOT altered with cmp/test instructions
- ▶ EFLAGS register IS changed by cmp/test to indicate less than, greater than, 0, etc.

EXAMPLES:

```
cmpl $1, %eax      # compare eax to 1
## EFLAGS bits set based on result of eax - 1
##  ZF (zero flag) now 1 if eax==1
##  SF (sign flag) now 1 if eax<1

testq %rax,%rax     # test rax against rax
## EFLAGS bits set based on result of rax & rax
##  ZF (zero flag) now 1 if rax==0 (falsey)
##  ZF (zero flag) now 0 if rax!=0 (truthy)
```

Jump Instruction Summary

- ▶ `jmp LAB`: **unconditional jump**, always go to another code location
- ▶ Control structures like `if/else/for/while` use `cmpX / testX` followed by **conditional jumps**
- ▶ `ja` used by compiler for `if(a < 0 || a > lim)`
Consider sign/unsigned to explain why
- ▶ `jmp *%rdx` allows **function pointers**, powerful but no time to discuss

<i>Instruction</i>	<i>Jump Condition</i>
<code>jmp LAB</code>	Unconditional jump
<code>je LAB</code>	Equal / zero
<code>jz LAB</code>	
<code>jne LAB</code>	Not equal / non-zero
<code>jnz LAB</code>	
<code>js LAB</code>	Negative ("signed")
<code>jns LAB</code>	Nonnegative
<code>jg LAB</code>	Greater-than signed
<code>jge LAB</code>	Greater-than-equal signed
<code>jl LAB</code>	Less-than signed
<code>jle LAB</code>	Less-than-equal signed
<code>ja LAB</code>	Above unsigned
<code>jae LAB</code>	Above-equal unsigned
<code>jb LAB</code>	Below unsigned
<code>jbe LAB</code>	Below-equal unsigned
<code>jmp *OPER</code>	Unconditional jump to variable address

Examine: Compiler Comparison Inversion

- ▶ Often compiler inverts comparisons
- ▶ $i < n$ becomes `cmpX / jge` (jump greater/equal)
- ▶ $i == 0$ becomes `cmpX / jne` (jump not equal)
- ▶ This allows “true” case to fall through immediately
- ▶ Depending on structure, may have additional jumps
 - ▶ `if(){ .. }` usually has a single jump
 - ▶ `if(){} else {}` may have a couple

```
## Assembly translation of
## if(rbx >= 2){
##   rdx = 10;
## }
## else{
##   rdx = 5;
## }
## return rdx;
    cmpq  $2,%rbx      # compare: rbx-0
    jl    .LESSTHAN    # goto less than
    ## if(rbx >= 2){
    movq  $10,%rdx     # greater/equal
    ## }
    jmp   .AFTER
.LESSTHAN:
    ## else{
    movq  $5,%rdx      # less than
    ## }
.AFTER:
    ## rdx is 10 if rbx >= 2
    ## rdx is 5 otherwise
    movq  %rdx,%rax
    ret
```

Exercise: Other Kinds of Conditions

Other Things to Look For

- ▶ `testl %eax,%eax` used to check zero/nonzero
- ▶ Followed by `je` / `jz` / `jne` / `jnz`
- ▶ Also works for NULL checks
- ▶ Negative Values, followed by `js` / `jns` (jump sign / jump no sign)

See `jmp_tests_asm.s`

- ▶ Trace the execution of this code
- ▶ Determine return value in `%eax`

cmov Family: Conditional Moves

- ▶ A family of instructions allows conditional movement of data into registers
- ▶ Can limit jumping in simple assignments

```
cmpq    %r8,%r9
cmovge  %r11,%r10  # if(r9 >= r8) { r10 = r11 }
cmovg   %r13,%r12  # if(r9 >  r8) { r12 = r13 }
```

- ▶ Note that condition flags are set on arithmetic operations
- ▶ cmpX is like subQ: both set FLAG bits the same
- ▶ Greater than is based on the SIGN flag indicating subtraction would be negative allowing the following:

```
subq    %r8,%r9    # r9 = r9 - r8
cmovge  %r11,%r10  # if(r9 >= 0) { r10 = r11 }
cmovg   %r13,%r12  # if(r9 >  0) { r12 = r13 }
```

Procedure Calls

Have seen basics so far:

```
call  PROCNAME  # call a function
        ## Pushes return address %rip onto stack adjusting

movl   $0,%eax   # set up return value
ret                    # return from function
        ## Pops old %rip off of stack adjusting %rsp
```

Need several additional notions

- ▶ Control Transfer?
- ▶ Where are arguments to functions?
- ▶ Return value?
- ▶ Anything special about the registers?
- ▶ How is the stack used?

Function/Procedure Control Transfer

call Instruction

1. Push the “caller” **Return Address** onto the stack
Return address is for instruction after call
2. Change rip to first instruction of the “callee” function

ret Instruction

1. Set rip to Return Address at top of stack
2. Pop the Return Address off the stack shrinking stack




Figure: Bryant/O'Hallaron Fig 3.26 demonstrates call/return in assembly

Example: Control Transfer with call

Example derived from `sum_range.s` file in code pack

BEFORE CALL

main: ...

```
0x555555554687 <+11>:      mov     $0x5,%esi
=> 0x55555555468c <+16>:      callq   0x55555555466a <sum_range>
0x555555554691 <+21>:      mov     %eax,%ebx
```

rip = 0x55555555468c -> call -> 0x555555554691

rsp = 0x7fffffff460

(gdb) stepi

AFTER CALL

sum_range:

```
=> 0x55555555466a <+0>: mov     $0x0,%eax
0x55555555466f <+5>: jmp     0x555555554676 <.TOP>
```

rip = 0x55555555466a

rsp = 0x7fffffff458 # pushed return address: rsp -= 8

(gdb) x/xg \$rsp

0x7fffffff458: 0x555555554691 # return address in main

Control Transfer with ret

BEFORE RET:

sum_range:...

0x555555554678 <+2>: jle 0x555555554671 <.BODY>

=> 0x55555555467a <+4>: repz retq

rip = 0x55555555467a -> return

rsp = 0x7fffffff458

(gdb) x/xg \$rsp

0x7fffffff458: 0x555555554691 # return address in main

(gdb) stepi

AFTER RET

0x555555554687 <+11>: mov \$0x5,%esi

0x55555555468c <+16>: callq 0x55555555466a <sum_range>

=> 0x555555554691 <+21>: mov %eax,%ebx

rip = 0x555555554691

rsp = 0x7fffffff460 # popped return address: rsp += 8

Stack Alignment

- ▶ According to the strict x86-64 ABI, must align `rsp` (stack pointer) to 16-byte boundaries when calling functions
- ▶ Will often see arbitrary pushes or subtractions to align
 - ▶ Always enter a function with old `rip` on the stack
 - ▶ Means that it is aligned to 8-byte boundary
- ▶ `rsp` changes must be undone prior to return

```
main:                                # enter with at 8-byte boundary
    subq    $8, %rsp                # align stack for func calls
    ...
    call    sum_range              # call function
    ...
    addq    $8, %rsp                # remove rsp change
    ret
```
- ▶ Failing to align the stack may work but may break
- ▶ Failing to “undo” stack pointer changes will likely result in return to the wrong spot : major problems

x86-64 Register/Procedure Convention

- ▶ Used by Linux/Mac/BSD/General Unix
- ▶ Params and return in registers if possible

Parameters and Return

- ▶ First 6 arguments are put into
 1. rdi / edi / di (arg 1)
 2. rsi / esi / si (arg 2)
 3. rdx / edx / dx (arg 3)
 4. rcx / ecx / cx (arg 4)
 5. r8 / r8d / r8w (arg 5)
 6. r9 / r9d / r9w (arg 6)
- ▶ Additional arguments are pushed onto the stack
- ▶ Return Value in rax / eax / ...

Caller/Callee Save

Caller save registers: alter freely

rax rcx rdx rdi rsi
r8 r9 r10 r11

Callee save registers: must restore these on return

rbx rbp r12 r13 r14
r15

Careful messing with stack pointer

rsp # stack pointer

Pushing and Popping the Stack

- ▶ If local variables are needed on the stack, can use `push` / `pop` for these
- ▶ `pushX %reg`: grow `rsp` (lower value), move value to top of main memory stack,
 - ▶ `pushq %rax`: grows `rsp` by 8, puts contents of `rax` at top
 - ▶ `pushl $25`: grows `rsp` by 4, puts constant 5 at top of stack
- ▶ `popX %reg`: move value from top of main memory stack to `reg`, shrink `rsp` (higher value)
 - ▶ `popl %eax`: move (`%rsp`) to `eax`, shrink `rsp` by 4

main:

```
    pushq    %rbp                # save register, aligns stack
                                   # like subq $8,%rsp; movq %rbp,(%rsp)
    call     sum_range           # call function
    movl     %eax, %ebp          # save answer
    ...
    call     sum_range           # call function, ebp not affected
    ...
    popq     %rbp                # restore rbp, shrinks stack
                                   # like movq (%rsp),%rbp; addq $8,%rsp
    ret
```

Exercise: Local Variables which need an Address

Compare code in files

- ▶ `swap_pointers.c` : familiar C code for swap via pointers
- ▶ `swap_pointers_asm.s` : hand-coded assembly version

Determine the following

1. Where are local C variables `x`, `y` stored in assembly version?
2. Where does the assembly version “grow” the stack?
3. Where does the assembly version “shrink” the stack?

Answers: Local Variables which need an Address

1. Where are local C variables x,y stored in assembly version?
2. Where does the assembly version “grow” the stack?

```
// C CODE
```

```
int x = 19, y = 31;
```

```
swap_ptr(&x, &y) // need main mem addresses for x,y
```

```
### ASSEMBLY CODE
```

```
main:                                # main() function
    subq    $8, %rsp                 # grow stack by 8 bytes
    movl    $19, (%rsp)              # move 19 to local variable x
    movl    $31, 4(%rsp)             # move 31 to local variable y
    movq    %rsp, %rdi               # load address of x into rdi
    leaq    4(%rsp), %rsi            # load address of y into rsi
    call    swap_ptr                 # call swap function
```

3. Where does the assembly version “shrink” the stack?

```
    addq    $8, %rsp                 # shrink stack by 8 bytes
    movl    $0, %eax                 # set return value
    ret
```

Diagram of Stack Arguments

- ▶ Compiler determines if local variables go on stack
- ▶ If so, calculates location as $\text{rsp} + \text{offsets}$

<pre>1 // C Code: locals.c 2 int set_buf(char *b, int *s); 3 int main(){ 4 // locals re-ordered on 5 // stack by compiler 6 int size = -1; 7 char buf[16]; 8 ... 9 int x = set_buf(buf, &size); 10 ... 11 }</pre>	<table border="1" style="border-collapse: collapse; width: 100%;"><thead><tr><th style="text-align: left; padding: 2px;">REG</th><th style="text-align: left; padding: 2px;">VALUE</th><th style="text-align: left; padding: 2px;">Name</th></tr></thead><tbody><tr><td style="padding: 2px;">rsp</td><td style="padding: 2px;">#1024</td><td style="padding: 2px;">top of stack during main</td></tr><tr><td style="padding: 2px;">MEM</td><td style="padding: 2px;"></td><td style="padding: 2px;"></td></tr><tr><td style="padding: 2px;">...</td><td style="padding: 2px;">...</td><td style="padding: 2px;">...</td></tr><tr><td style="padding: 2px;">#1031</td><td style="padding: 2px;">h</td><td style="padding: 2px;">buf[3]</td></tr><tr><td style="padding: 2px;">#1030</td><td style="padding: 2px;">s</td><td style="padding: 2px;">buf[2]</td></tr><tr><td style="padding: 2px;">#1029</td><td style="padding: 2px;">u</td><td style="padding: 2px;">buf[1]</td></tr><tr><td style="padding: 2px;">#1028</td><td style="padding: 2px;">p</td><td style="padding: 2px;">buf[0]</td></tr><tr><td style="padding: 2px;">#1024</td><td style="padding: 2px;">-1</td><td style="padding: 2px;">size</td></tr></tbody></table>	REG	VALUE	Name	rsp	#1024	top of stack during main	MEM			#1031	h	buf[3]	#1030	s	buf[2]	#1029	u	buf[1]	#1028	p	buf[0]	#1024	-1	size
REG	VALUE	Name																										
rsp	#1024	top of stack during main																										
MEM																												
...																										
#1031	h	buf[3]																										
#1030	s	buf[2]																										
#1029	u	buf[1]																										
#1028	p	buf[0]																										
#1024	-1	size																										
<pre>1 ## EQUIVALENT ASSEMBLY 2 main: 3 subq \$24, %rsp # space for buf/size and stack alignment 4 movl \$-1, (%rsp) # old rip already in stack so: 20+4+8 = 32 5 # initialize buf and size: main line 6 6 leaq 0(%rsp), %rdi # address of size arg1 7 leaq 4(%rsp), %rsi # address of buf arg2 8 call set_buf # call function, aligned to 16-byte boundary 9 movl %eax,%r8 # get return value 10 ... 11 addq \$24, %rsp # shrink stack size</pre>																												

Historical Aside: Base Pointer rbp was Important

```
int bar(int, int, int);
int foo(void) {
    int x = callee(1, 2, 3);
    return x+5;
}
```

- ▶ 32-bit x86 / IA32 assembly used rbp as bottom of stack frame, rsp as top.
- ▶ Push all arguments onto the stack when calling changing both rsp and rbp
- ▶ x86-64: default rbp to general purpose register, not used for stack tracking

Old x86 / IA32 calling sequence: set both %esp and %ebp for function call foo:

```
    pushl %ebp                # modifying ebp, save it
    ## Set up for function call to bar()
    movl  %esp,%ebp          # new frame for next function
    pushl 3                   # push all arguments to
    pushl 2                   # function onto stack
    pushl 1                   # no regs used
    call bar                  # call function, return val in %eax
    ## Tear down for function call bar()
    movl  %ebp,%esp           # restore stack top: args popped
    ## Continue with function foo()
    addl 5,%eax               # add onto answer
    popl  %ebp               # restore previous base pointer
    ret
```